



UNIVERSITY OF  
**WATERLOO**

# CS 456/656 Computer Networks

## Lecture 18: Router/Switch Architecture


Mina Tahmasbi Arashloo and Uzma Maroof  
Fall 2025

# A note on the slides

Adapted from the slides that accompany this book.

All material copyright 1996-2023  
J.F Kurose and K.W. Ross, All Rights Reserved


And lecture notes from Anirudh Sivaraman, NYU



*Computer Networking: A Top-Down Approach*  
8<sup>th</sup> edition  
Jim Kurose, Keith Ross  
Pearson, 2020


# What we discussed before

- Packets can be buffered in routers
  - delay and loss
  - network congestion




# What we discussed before

- Packets can be buffered in routers
  - delay and loss
  - network congestion



# What we discussed before

- Packets can be buffered in routers
  - delay and loss
  - network congestion




This Lecture: How are packets buffered and managed in routers?

# Router architecture and buffer management


- Router architecture
  - Output queueing
  - Input queueing
  - ...
- Buffer management and scheduling

# What is inside a router?

high-level view of generic router architecture:




# Input port functions




# Switching fabrics

- transfer packets from input links to appropriate output links
- Suppose
  - All packets are of the same size
  - Define the time it takes to send/receive a packet on a port as our time unit, and call it a *tick* (today, that's usually a few nanoseconds!)



# Switching fabrics


- In each tick, we can
  - receive at most a packet on each input port (up to N ports)
  - send at most a packet on each output port (up to N ports)



# Switching fabrics

On each tick, we can  
receive a packet on each input port  
send a packet on each output port


- Ideally, the switching fabric can move  $N$  packets in each tick
  - If the link rates are  $R$ , an ideal switching fabric moves packets at rate  $NR$ .



# Switching fabrics

On each tick, we can  
receive a packet on each input port  
send a packet on each output port


- If two or more input ports have a packet destined to the same output port
  - only one can go out in the next tick(s)
  - the rest have to wait somewhere



# Switching fabrics

On each tick, we can  
receive a packet on each input port  
send a packet on each output port

- If two or more input ports have a packet destined to the same output port
  - only one can go out in the next tick(s)
  - the rest have to wait somewhere




# Router architecture and buffer management

- Router architecture
  - Output queueing
  - Input queueing
  - ...
- Buffer management and scheduling

# Output queuing

- N separate memories, one for each output port
- Input port receives a packet, and then puts it in the memory of the output port it is supposed to exit from
- Output port pulls the next packet from its corresponding memory



# Router architecture and buffer management

- Router architecture

- Output queueing
- Input queueing
- ...


- Buffer management and scheduling

# Input queuing

- N separate memory pools, one for each input port.
- Input port receives a packet, puts it in the dedicated memory for that input port.
  - 1 enqueue per tick
- Output port gets the next packet from the memory of one of the input ports that have packets destined to it
  - 1 dequeue per tick



# Input queuing

At most one enqueue in each memory per tick





# Input queuing

At most one dequeue from each memory per tick




# Input queuing

At most one dequeue from each memory per tick



# Input queuing


How do we coordinate packets between inputs and outputs?



A single buffer per input can lead to head of line blocking... look up virtual output queuing for a common solution

In every tick

- Each input sends to at most one output
- Each output receives from at most each input



Coordinates between inputs and outputs by solving a bi-partite matching problem!

# Architecture trends

- Today, there is a renewed interest in output-queued (and shared memory) architectures
- Data centers have *many* switches (100s of thousands)
- To keep the costs down, vendors have reduced the amount of memory available for buffering in these switches
  - Easier, e.g., compared to a WAN, to keep the queues shorter in DCs, specially with the help of congestion control algorithms.
- Easier to make smaller high-speed memory with multiple enqueues and/or dequeues per tick
- With output-queued (or shared-memory architectures), no need for dealing with efficient scheduling of a crossbar.

# Router architecture and buffer management

- Router architecture

- Output queueing
- Input queueing
- ...

- Buffer management and scheduling

# Queue/Buffer management and scheduling

- Independent of where the queues are in the router architecture, there are some important questions:
  - *Buffer size*: How large should a buffer be?
  - *Queue management*: When the queue is full, which packet do we drop? What do we do when the queue starts building up?
  - *Packet scheduling*: Which packet in the queue gets dequeued first? Should it be first-in first-out? Something else?


# How large should a buffer be?

*There is no easy answer:*

- Too small: can't absorb bursts, keeps dropping packets
- Too large: can hurt performance
  - *buffer bloat*: When the buffer is too large, it will take a long time to fill up before a packet is dropped (however, TCP only realizes there is congestion when a packet is dropped and will not decrease its sending rate). In the meantime, all packets will experience increasing queueing delay.
  - Delay-based congestion control algorithms do better here.

# Queue management – Drop policy

- When a new packet arrives to a full queue, which packet do we drop?
- *Tail drop*: drop arriving packet
- *Priority*: drop/remove based on priority
  - E.g., if the incoming packet has higher priority than a packet already in the queue, drop the lower priority packet and insert the incoming packet into the queue.




# Queue management – Marking

- When the queue starts filling up, one strategy is to mark packets to signal the onset of congestion to the end points
- Recall our discussion about Early Congestion Notification (ECN) and its role in congestion control
- When should we start/stop marking packets?
- Which packets do we mark?
  - All packets after the queue size passes a threshold?
  - From the flow with the most packets in the queue?
  - ...


# Packet scheduling

- So far, we have assumed that our queues are *first in first out (FIFO)*



- But, there are other packet scheduling algorithms as well.

# Packet scheduling



# Additional Slides